

CS201 – Lecture 3

Linking

RAOUL RIVAS

PORTLAND STATE UNIVERSITY

A solid green horizontal bar spanning the width of the slide at the bottom.

Announcements

Multiple Source files

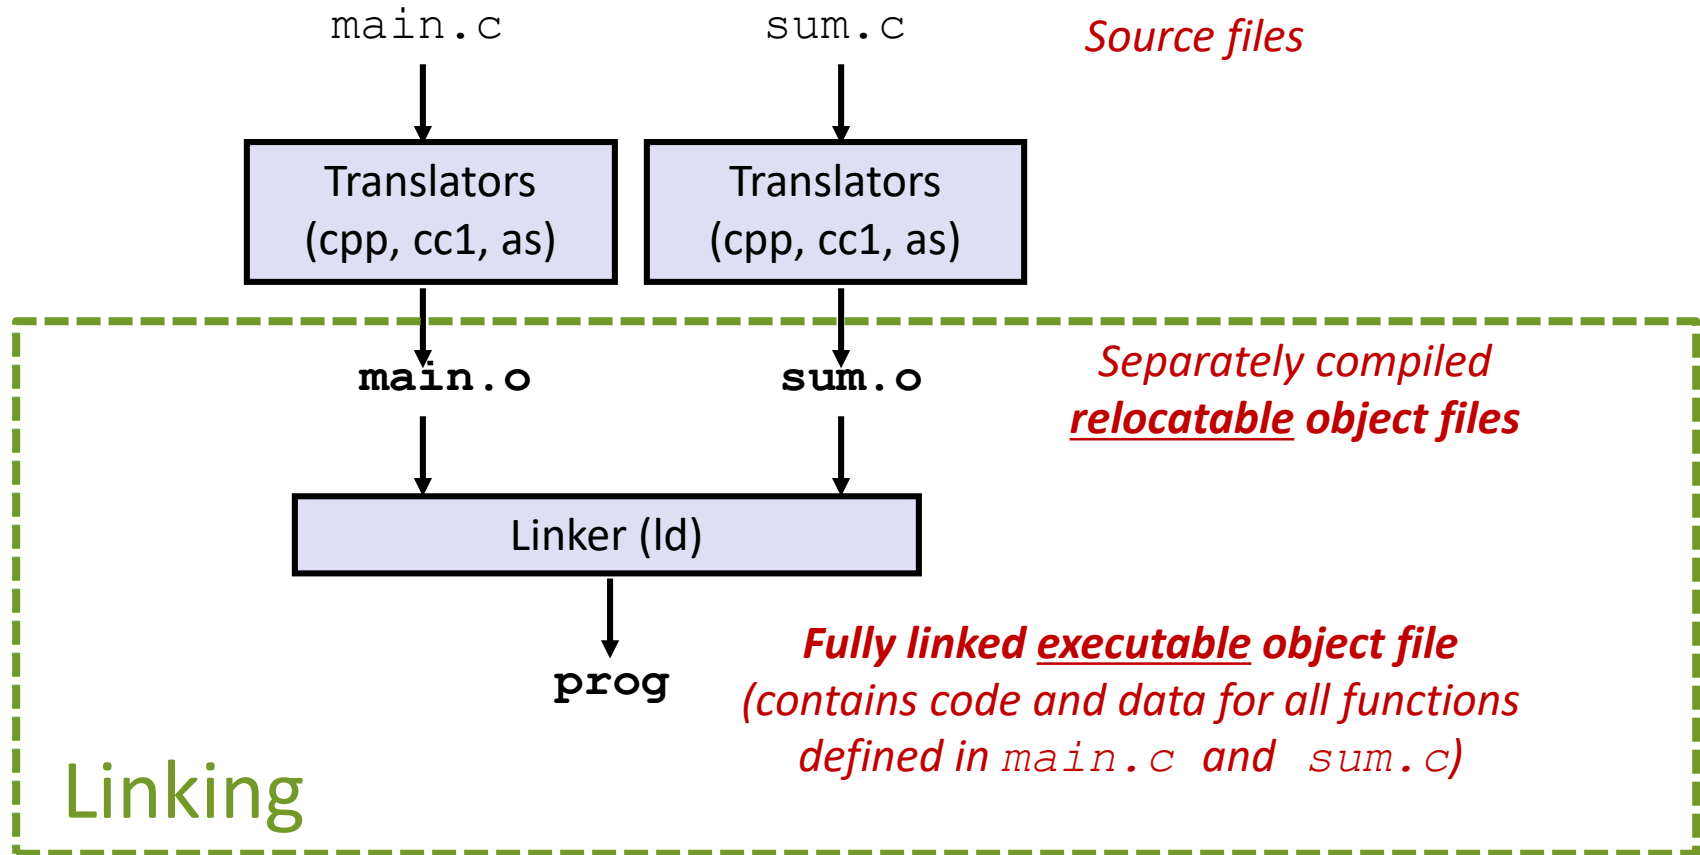
main.c

```
int sum(int *a, int n);  
  
int array[2] = {1, 2};  
  
int main()  
{  
    int val = sum(array, 2);  
    return val;  
}
```

sum.c

```
int sum(int *a, int n)  
{  
    int i, s = 0;  
  
    for (i = 0; i < n; i++) {  
        s += a[i];  
    }  
    return s;  
}
```

Why Linkers?



Why Linkers?

- Reason 1: Modularity
 - Applications can be very large
 - Linux: 15 million lines of code
 - Windows: 50 million lines of code
 - Office: 40 million lines of code
 - Program can be written as a collection of smaller source files, rather than one monolithic mass.
 - Can build libraries of common functions (more on this later)
 - e.g., Math library, standard C library

Why Linkers?


- Reason 2: Efficiency
 - Time: Separate compilation
 - Change one source file, compile, and then relink.
 - No need to recompile other source files.
 - Space: Libraries
 - Common functions can be aggregated into a single file...
 - Yet executable files and running memory images contain only code for the functions they actually use.

Object Files

- Relocatable object file (`.o` file)
 - Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each `.o` file is produced from exactly one source (`.c`) file
- Executable object file (`a.out` file)
 - Contains code and data in a form that can be copied directly into memory and then executed.
 - These are `*.EXE` and `*.COM` files in Windows
 - **Non Relocatable!**
- Shared object file (`.so` file)
 - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
 - Called *Dynamic Link Libraries* (DLLs) by Windows

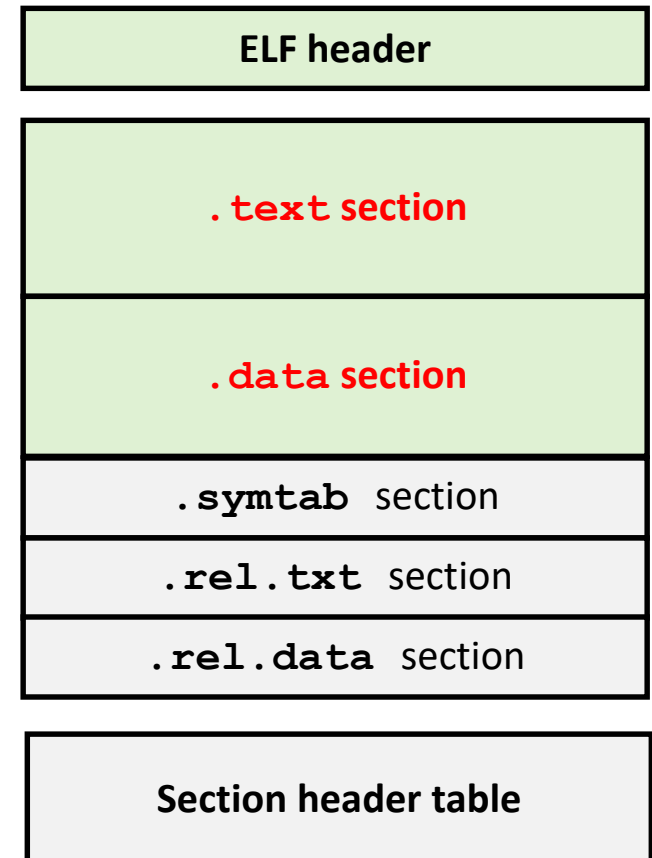
Executable and Linkable Format (ELF)

- Standard binary format for object files in Unix and Linux
- One unified format for
 - Relocatable object files (`.o`),
 - Executable object files (`a.out`)
 - Shared object files (`.so`)
- Generic name: ELF binaries
- Modular design
 - Program header table contains pointers to each section



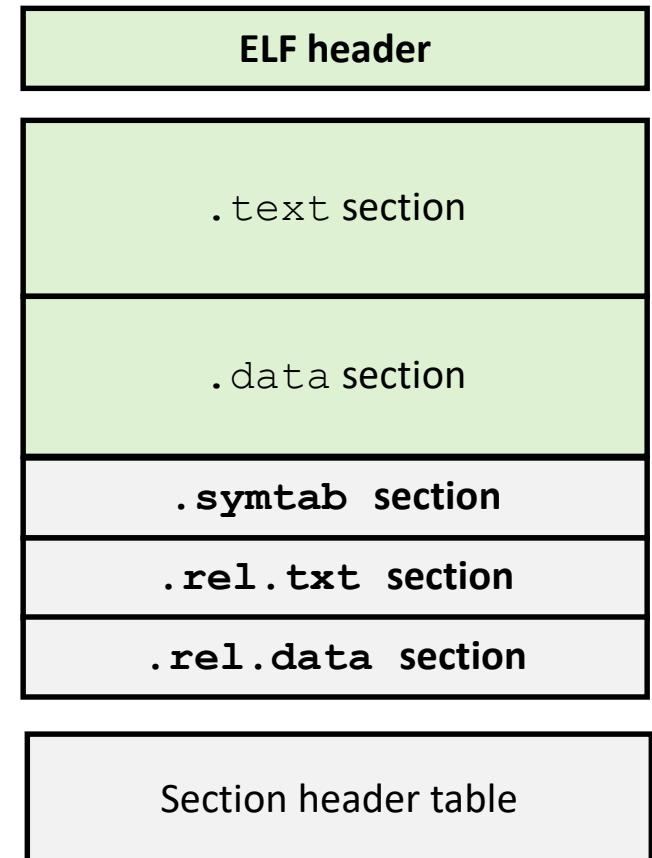
ELF Structure

- Elf header
 - Word size, byte ordering, file type (.o, exec, .so), machine type, etc.
- **.text section**
 - Code
- **.data section**
 - Initialized global variables
- .symtab section
 - Symbol table
 - Procedure and static variable names
 - Section names and locations
- Section header table
 - Offsets and sizes of each section for Relocatable Object Files



ELF Structure – Relocation Structures

- Relocation data is stored in the
 - `.rel.text` section
 - Relocation info for `.text` section
 - Addresses of instructions that will need to be modified in the executable
 - Instructions for modifying.
 - `.rel.data` section
 - Relocation info for `.data` section
 - Addresses of pointer data that will need to be modified in the merged executable



0

Step 1: Symbol Resolution

- Step 1: Symbol resolution

- Programs define and reference *symbols* (global variables and functions):

- `void swap() {...}` `/* define symbol swap */`
- `swap();` `/* reference symbol swap */`
- `int *xp = &x;` `/* define symbol xp,`
 `reference x */`

- Symbol definitions are stored in object file (by assembler) in *symbol table*.

- Symbol table is an array of `structs`
- Each entry includes name, size, and location of symbol.

- **During symbol resolution step, the linker associates each symbol reference with exactly one symbol definition.**

Linker Symbols

- Global symbols
 - Symbols defined by module m that can be referenced by other modules.
 - E.g.: non-**static** C functions and non-**static** global variables.
- External symbols
 - Global symbols that are referenced by module m but defined by some other module.
- Local symbols
 - Symbols that are defined and referenced exclusively by module m .
 - E.g.: C functions and global variables defined with the **static** attribute.
 - **Local linker symbols are *not* local program variables**

Step 1: Symbol Resolution

Reference to a Local Symbol

Local Symbol

main.c

```
int sum(int *a, int n);  
  
int array[2] = {1, 2};  
  
int main()  
{  
    int val = sum(array, 2);  
    return val;  
}
```

Definition of
a Global
Symbol

Linker knows
nothing of `val`

Reference to an
External Symbol

sum.c

```
int sum(int *a, int n)  
{  
    int i, s = 0;  
    for (i = 0; i < n; i++) {  
        s += a[i];  
    }  
    return s;  
}
```

Definition of a
Global symbol

Linker knows
nothing of `i` or `s`

Local Symbols

- Local non-static C variables vs. local static C variables
 - local non-static C variables: stored on the stack
 - local static C variables: stored in `.data`

```
int f()
{
    static int x = 0;
    return x;
}


int g()
{
    static int x = 1;
    return x;
}
```

Compiler allocates space in `.data` for each definition of `x`


Creates local symbols in the symbol table with unique names, e.g., `x.f` and `x.g`.

Step 2: Relocation


Relocatable Object Files



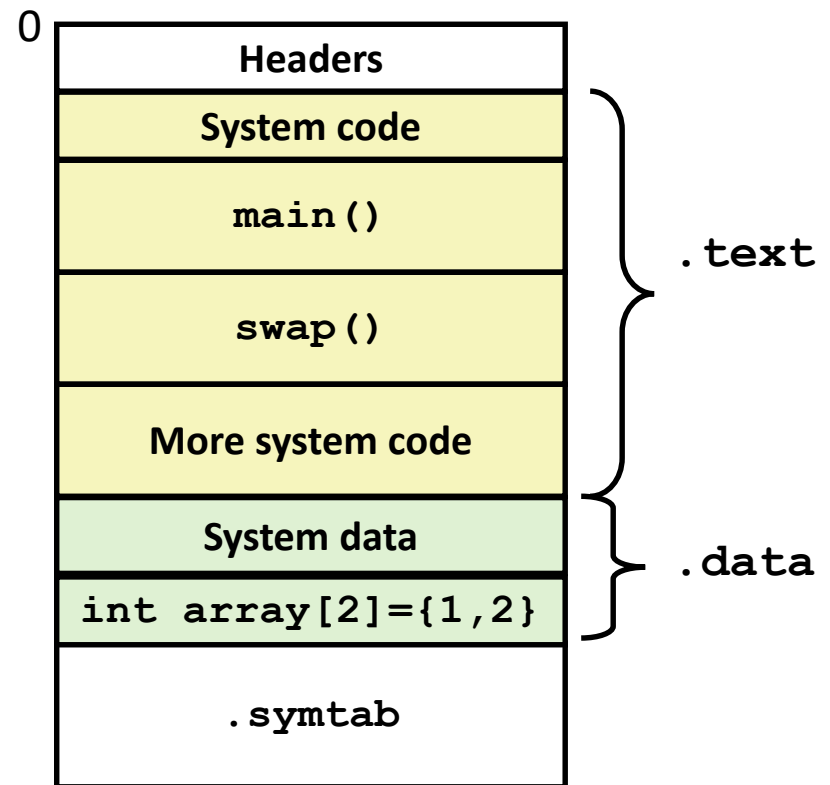
main.o



sum.o



Executable Object File



Relocation Entries

```
int array[2] = {1, 2};

int main()
{
    int val = sum(array, 2);
    return val;
}
```

main.c

0000000000000000 <main>:

0:	48 83 ec 08	sub	\$0x8,%rsp	
4:	be 02 00 00 00	mov	\$0x2,%esi	
9:	bf 00 00 00 00	mov	\$0x0,%edi	# %edi = &array
		a: R_X86_64_32	array	# Relocation entry
e:	e8 00 00 00 00	callq	13 <main+0x13>	# sum()
		f: R_X86_64_PC32	sum-0x4	# Relocation entry
13:	48 83 c4 08	add	\$0x8,%rsp	
17:	c3	retq		

main.o

Relocated .text section

00000000004004d0 <main>:

4004d0:	48 83 ec 08	sub	\$0x8,%rsp
4004d4:	be 02 00 00 00	mov	\$0x2,%esi
4004d9:	bf 18 10 60 00	mov	\$0x601018,%edi # %edi = &array
4004de:	e8 05 00 00 00	callq	4004e8 <sum> # sum()
4004e3:	48 83 c4 08	add	\$0x8,%rsp
4004e7:	c3	retq	

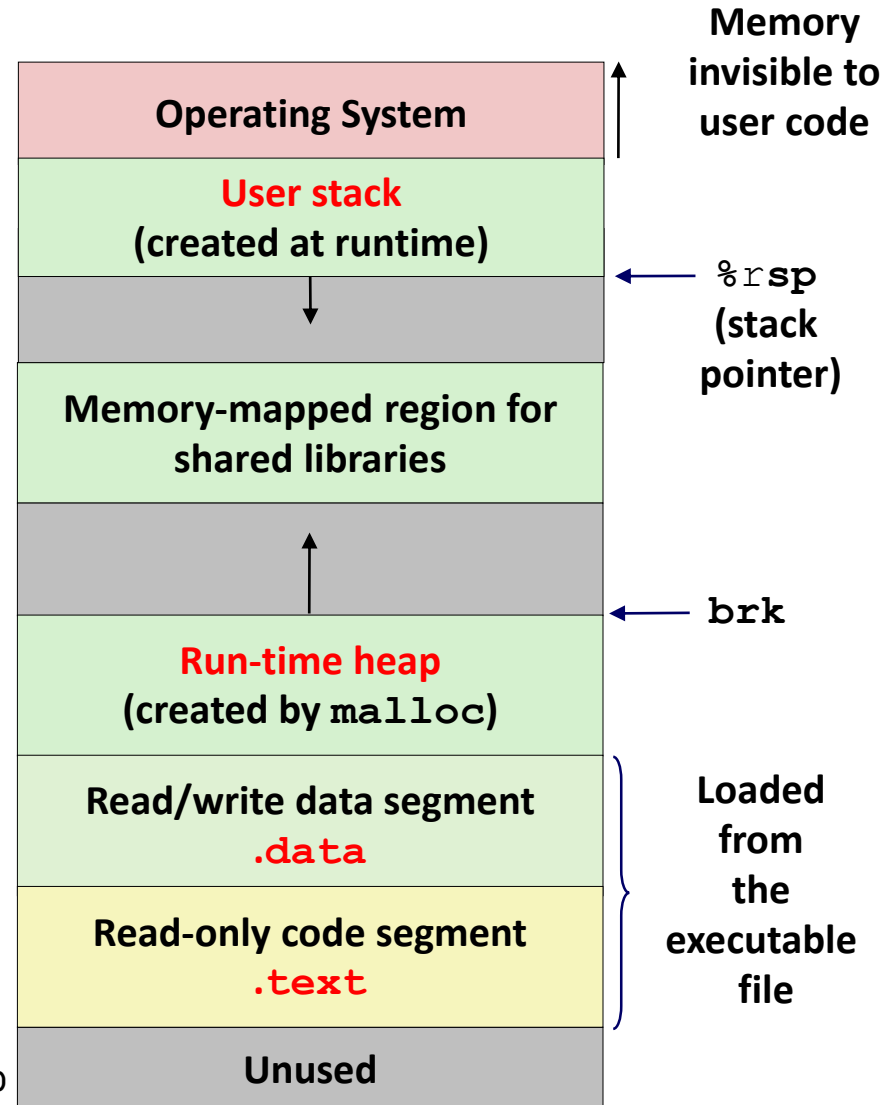

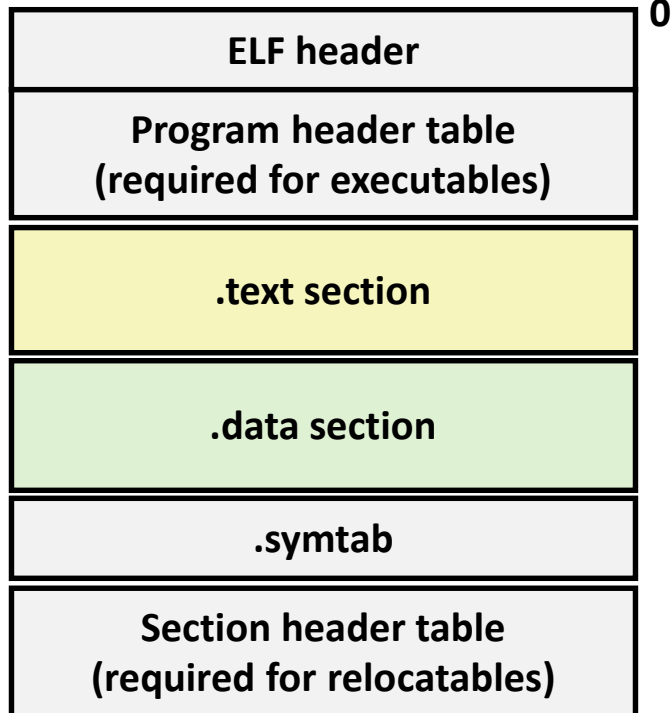
00000000004004e8 <sum>:

4004e8:	b8 00 00 00 00	mov	\$0x0,%eax
4004ed:	ba 00 00 00 00	mov	\$0x0,%edx
4004f2:	eb 09	jmp	4004fd <sum+0x15>
4004f4:	48 63 ca	movslq	%edx,%rcx
4004f7:	03 04 8f	add	(%rdi,%rcx,4),%eax
4004fa:	83 c2 01	add	\$0x1,%edx
4004fd:	39 f2	cmp	%esi,%edx
4004ff:	7c f3	j1	4004f4 <sum+0xc>
400501:	f3 c3	repz retq	

Using PC-relative addressing for sum(): $0x4004e8 = 0x4004e3 + 0x5$

Loading an Executable

Executable Object File



Loading an Executable

- User Stack
 - Allocate local variables
 - Store the state of the registers of the caller including the return address of the caller
- Run-Time Heap
 - Used to allocate dynamic memory (malloc)
- Read Write Segment
 - Store statically allocated globals (.data segment)
- Read-Only Segment
 - Store program code (.text segment)

Memory Allocation

```
char big array[1L<<24]; /* 16 MB */


int global = 0;

int useless() { return 0; }

int main ()
{
    void *p1,*p4;
    int local = 0;

    p1 = malloc(1L << 28); /* 256 MB */
    p4 = malloc(1L << 8); /* 256 B */

    /* Some print statements ... */
}
```

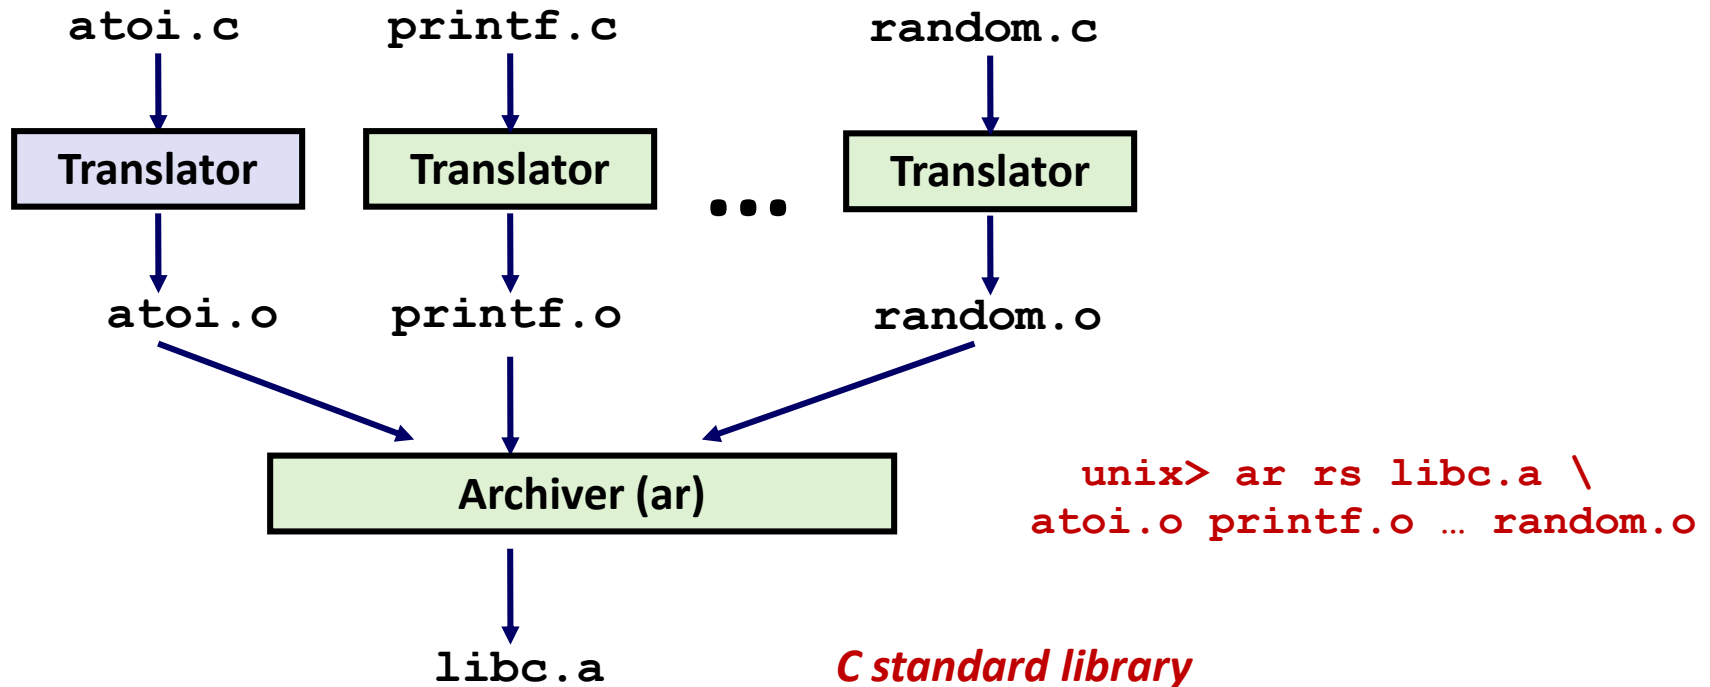


Where does everything go?

Static Libraries

- **Static libraries** (.a archive files)
 - Concatenate related relocatable object files into a single file with an index (called an *archive*).
 - Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives.
 - If an archive member file resolves reference, link it into the executable.

Static Libraries



- Archiver allows incremental updates
- Recompile function that changes and replace .o file in archive.

The C library

`libc.a` (the C standard library)

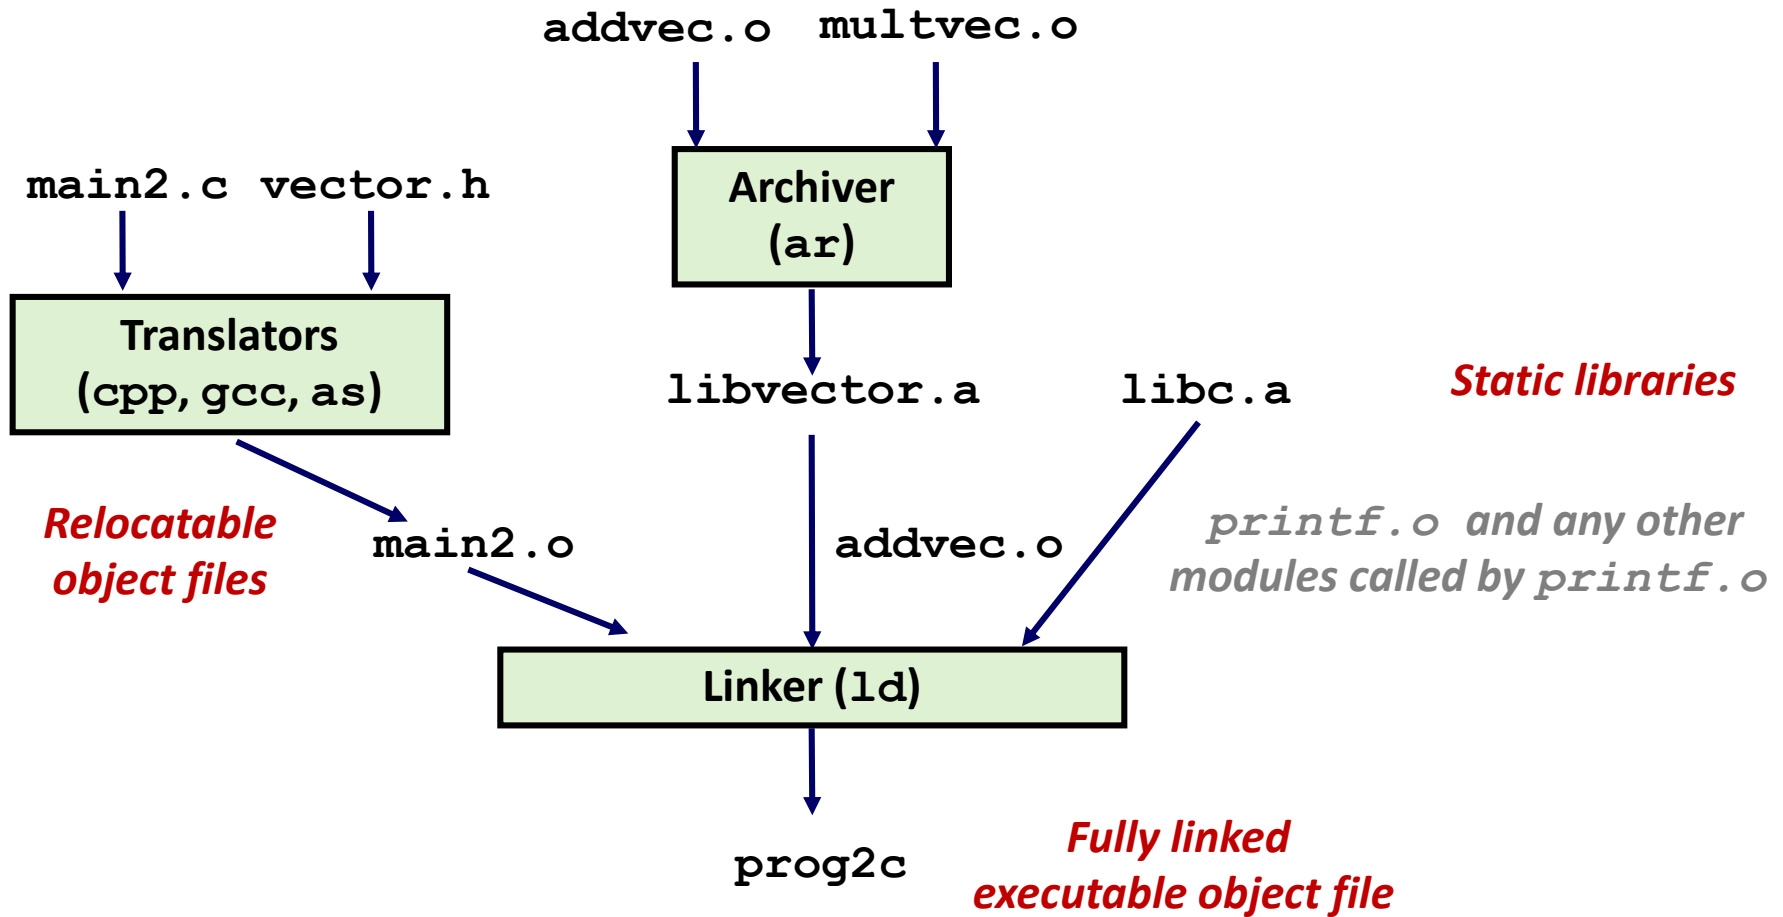
- 4.6 MB archive of 1496 object files.
- I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math

`libm.a` (the C math library)

- 2 MB archive of 444 object files.
- floating point math (sin, cos, tan, log, exp, sqrt, ...)

```
% ar -t libc.a | sort
fork.o
fprintf.o
fpu_control.o
fputc.o
freopen.o
fscanf.o
fseek.o
fstab.o
```

Linking Static Libraries



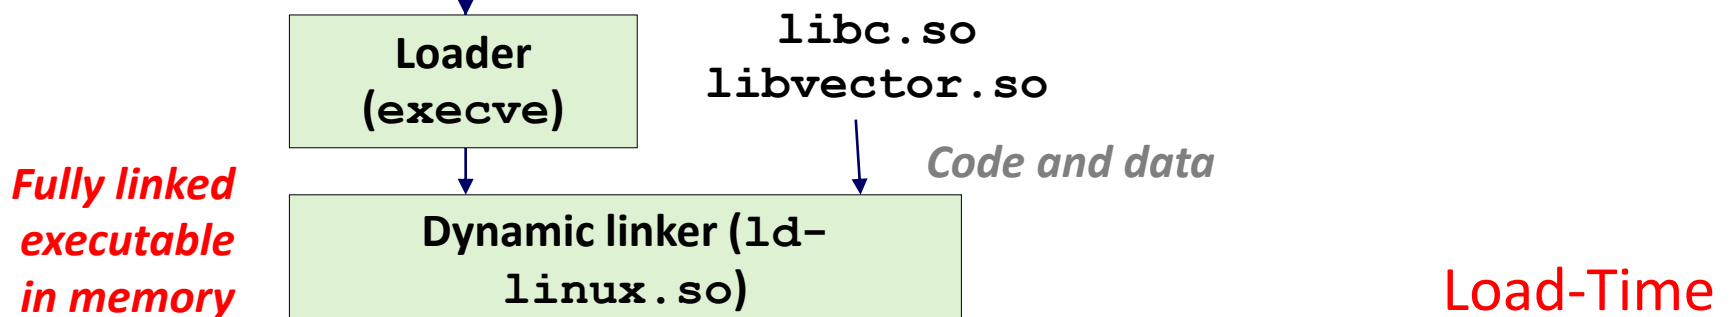
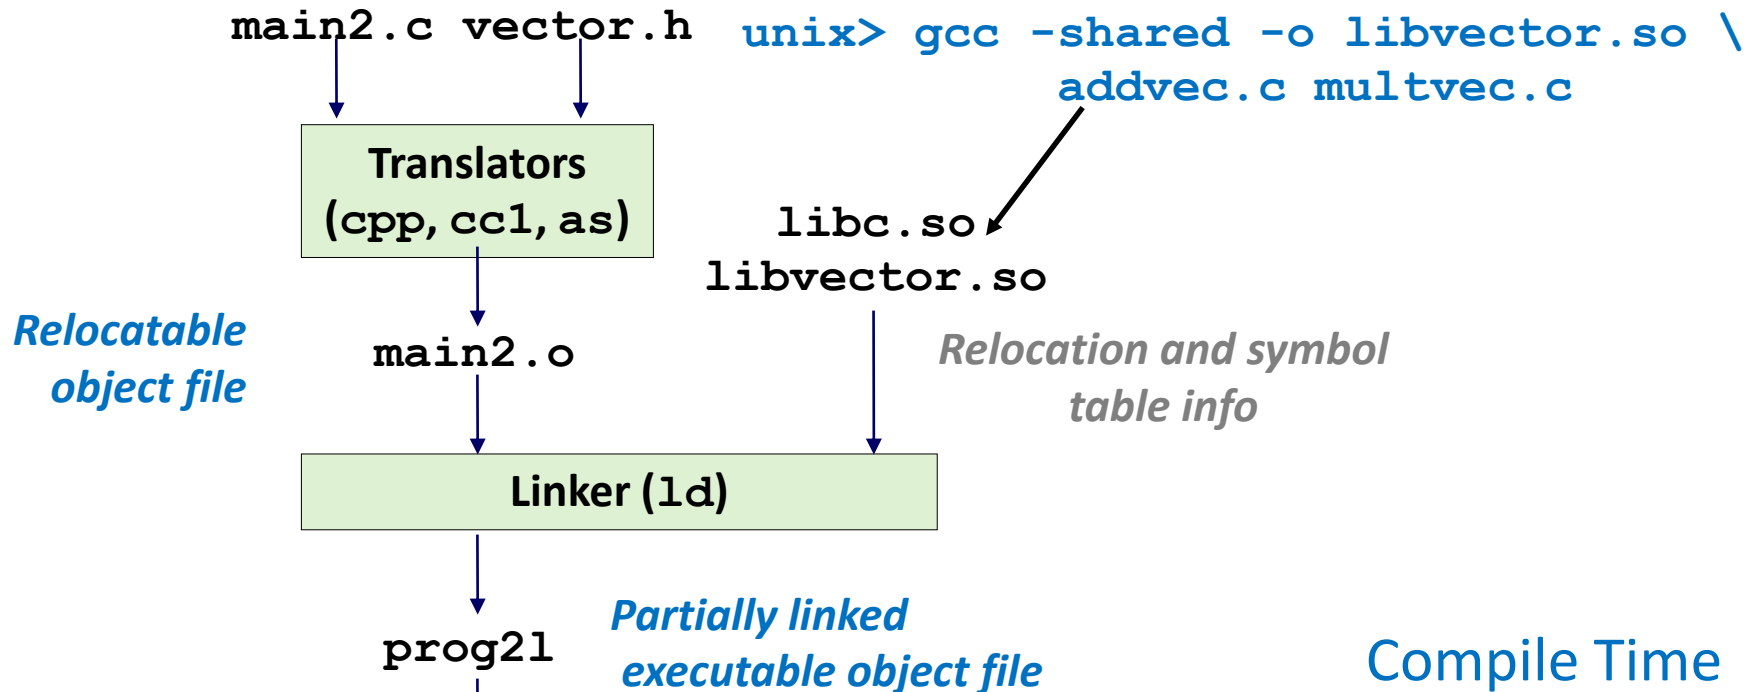
Shared Libraries

- Static libraries have the following disadvantages:
 - Duplication in the stored executables (every function needs libc)
 - Duplication in the running executables
 - Minor bug fixes of system libraries require each application to explicitly relink
- **Modern solution: Shared Libraries** Object files that contain code and data that are loaded and linked into an application *dynamically*, at either *load-time* or *run-time*
 - Also called: dynamic link libraries, DLLs, `.so` files

Dynamic Linking

- Dynamic linking can occur when executable is first loaded and run (load-time linking).
 - Common case for Linux, handled automatically by the dynamic linker (**ld-linux.so**) .
 - Standard C library (**libc.so**) usually dynamically linked.
- Dynamic linking can also occur after program has begun (run-time linking).
 - In Linux, this is done by calls to the **dlopen()** interface .
 - Add-ins and Plug-ins
 - High-performance web servers.
 - Runtime library inter-positioning.

Dynamic Linking at Load-time



Dynamic Linking at Run-time

```
int main()
{
    void *handle;
    void (*addvec)(int *, int *, int *, int);


    /* Dynamically load the shared library that contains
    addvec() */
    handle = dlopen("./libvector.so", RTLD_LAZY);
    /* Get a pointer to the addvec() function we just loaded */
    addvec = dlsym(handle, "addvec");
    /* Now we can call addvec() just like any other function */
    addvec(x, y, z, 2);
    /* Unload the shared library */
    dlclose(handle);
    return 0;
}
```

Library Linking Timeline


- Compile-Time Linking (Static)



- Load-Time Linking (Dynamic)



- Run-Time Linking (Dynamic)



Summary

- Linking is a technique that allows programs to be constructed from multiple object files.
- Linking can happen at different times in a program's lifetime:
 - Compile time (when a program is compiled)
 - Load time (when a program is loaded into memory)
 - Run time (while a program is executing)
- Understanding linking can help you avoid nasty errors and make you a better programmer.
- Read <http://www.ibm.com/developerworks/library/l-dynamic-libraries/> to learn more about Dynamic Linking in Linux